



# Intelligent Document Reader

Upload PDF



Yarlapa...vityAI.pdf  
6.1KB



## Extracted Text

SUKADEV YARLAPALLI AI / ML Engineer — LLMs, Agentic AI & MLOps Hyderabad, India | +91-6301008324 | [yarlapallisukadev@gmail.com](mailto:yarlapallisukadev@gmail.com) | [linkedin.com/in/sukadev382](https://linkedin.com/in/sukadev382) | GitHub PROFESSIONAL SUMMARY Builder-minded AI/ML engineer focused on shipping LLM-powered products end-to-end. Hands-on with LangChain, agentic workflows (planning, tool use, multi-step reasoning), recommendation systems, and MLOps (CI/CD, deployment, monitoring). Comfortable taking an idea from zero to a shipped product, iterating fast, and owning outcomes in a startup environment. TECHNICAL SKILLS (cid:127) Languages: Python, SQL, JSON, XML (cid:127) LLMs & Agents: LangChain, LangGraph, OpenAI / Gemini / Groq APIs, RAG, Function/Tool Calling, Prompt Engineering, Vector DBs (FAISS, Chroma, Pinecone) (cid:127) Agentic AI: ReAct, Planner-Executor, Multi-step Reasoning, Tool Use, Memory, Autonomous Workflows (cid:127) ML / Data: Scikit-Learn, Pandas, NumPy, NLP, OCR, Recommendation Systems (Collaborative + Content-based, Hybrid), Feature Engineering (cid:127) MLOps & Deployment: Docker, GitHub Actions (CI/CD), FastAPI, Streamlit, Hugging Face Spaces, Model Monitoring, Logging (cid:127) Automation & BI: Power Automate, PowerApps, Power BI, Databricks, REST APIs PROJECTS Agentic Research Assistant (LLM + LangChain Agents) Python, LangChain, LangGraph, OpenAI, FAISS, FastAPI, Docker (cid:127) Built a multi-step agentic AI that takes a single prompt, plans sub-tasks, picks tools (web search, RAG over PDFs, Python REPL, SQL), and ships a finished research brief without human hand-holding. (cid:127) Implemented a Planner-Executor loop with LangGraph + ReAct reasoning, tool-calling, and short/long-term memory backed by FAISS vector store. (cid:127) Wrapped the agent in a FastAPI service, containerized with Docker, and deployed with GitHub Actions CI/CD — average task completion in under 40 seconds end-to-end. Personalized Content Recommendation Engine Python, Scikit-Learn, Implicit, FastAPI, Streamlit (cid:127) Designed a hybrid recommendation engine combining collaborative filtering (ALS) and content-based embeddings to predict what users actually want next. (cid:127) Boosted offline top-10 recall by 28% vs. a popularity baseline on a 100K+ interaction dataset; added a cold-start fallback using LLM-generated item embeddings. (cid:127) Exposed real-time recommendations through a FastAPI endpoint with sub-100ms latency and an interactive Streamlit demo. LLM-Powered Workflow Automation Agent Python, LangChain, OpenAI Functions, REST APIs, Power Automate (cid:127) Built an agent that converts plain-English requests ("summarize yesterday's tickets and email the team") into executed multi-step work across Gmail, Sheets, and SharePoint. (cid:127) Used OpenAI function-calling + LangChain tools to safely orchestrate APIs with retries, validation, and audit logs — eliminated 15+ hours/week of manual ops work. (cid:127) Added guardrails: structured output schemas, action confirmation, and per-tool rate limiting. MLOps Pipeline for Continuous Model Delivery Python, Docker, GitHub Actions, FastAPI, Hugging

Face, Prometheus-style logging (cid:127) Stood up a full MLOps pipeline: data validation fi training fi evaluation fi containerized deploy fi live monitoring, triggered automatically on every git push. (cid:127) Used GitHub Actions for CI/CD, Docker for reproducible builds, and a lightweight monitoring layer to track drift, latency, and error rates in production. (cid:127) Cut model release cycle from days to under 15 minutes, enabling ship-fast / iterate-fast experimentation. Intelligent Document Reader (NLP + OCR) Python, Tesseract OCR, NLP, Regex, LangChain (cid:127) Developed an OCR + NLP pipeline that extracts, classifies, and summarizes scanned invoices and PDFs with 94% extraction accuracy. (cid:127) Layered an LLM on top to answer natural-language questions over the extracted content (RAG-style). (cid:127) Reduced document processing time from ~6 minutes to under 30 seconds per file. EDUCATION Imarticus Learning, Hyderabad — Post-Graduation in Data Science & Analytics (Advanced ML Track), 2025–2026 Oxford Degree College, Sri Venkateswara University — B.Sc. (Maths, Physics, Chemistry), 2022–2025 | CGPA: 6.5 CERTIFICATIONS (cid:127) Post-Graduation Certification in Data Science & Analytics with Advanced ML Track — Imarticus Learning (cid:127) Python for Data Science & Machine Learning — Self-paced, project-based learning ACHIEVEMENTS (cid:127) Shipped 5+ end-to-end AI/ML projects on GitHub & Hugging Face, including LLM agents and recommendation systems. (cid:127) Took multiple ideas from zero to deployed product without supervision — owning design, build, and release.

Ask a Question

what are my skills?

Get Answer

## Answer

**TECHNICAL SKILLS** (cid:127) Languages: Python, SQL, JSON, XML (cid:127) LLMs & Agents: LangChain, LangGraph, OpenAI / Gemini / Groq APIs, RAG, Function/Tool Calling, Prompt Engineering, Vector DBs (FAISS, Chroma, Pinecone) (cid:127) Agentic AI: ReAct, Planner-Executor, Multi-step Reasoning, Tool Use, Memory, Autonomous Workflows (cid:127) ML / Data: Scikit-Learn, Pandas, NumPy, NLP, OCR, Recommendation Systems (Collaborative + Content-based, Hybrid), Feature Engineering (cid:127) MLOps & Deployment: Docker, GitHub Actions (CI/CD), FastAPI, Streamlit, Hugging Face Spaces, Model Monitoring, Logging (cid:127) Automation & BI: Power Automate, PowerApps, Power BI, Databricks, REST APIs **PROJECTS** Agentic Research Assistant (LLM + LangChain Agents) Python, LangChain, LangGraph, OpenAI, FAISS, FastA